

```
#Importing the required libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
df = pd.read_csv("uber.csv")
```

pre process the dataset

```
df.head()
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354		
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225		
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770		
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844		
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085		



```
df.info() #To get the required information of the dataset
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 187717 entries, 0 to 187716
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0             187717 non-null  int64
1   key                    187717 non-null  object
2   fare_amount            187717 non-null  float64
3   pickup_datetime        187717 non-null  object
4   pickup_longitude       187717 non-null  float64
5   pickup_latitude        187717 non-null  float64
6   dropoff_longitude      187716 non-null  float64
7   dropoff_latitude       187715 non-null  float64
8   passenger_count        187716 non-null  float64
dtypes: float64(6), int64(1), object(2)
memory usage: 12.9+ MB

```

df.columns #TO get number of columns in the dataset

```

Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
      'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
      'dropoff_latitude', 'passenger_count'],
      dtype='object')

```

df = df.drop(['Unnamed: 0', 'key'], axis= 1) #To drop unnamed column as it isn't required

df.head()

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.723217	1
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.750325	1
2	10.0	2009-08-24	-73.995010	40.740770	-73.995000	40.750000	1

df.shape #To get the total (Rows,Columns)

```
(187717, 7)
```

df.dtypes #To get the type of each column

```
fare_amount      float64
pickup_datetime  object
pickup_longitude  float64
pickup_latitude   float64
dropoff_longitude float64
dropoff_latitude  float64
passenger_count   float64
dtype: object
```

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 187717 entries, 0 to 187716
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount           187717 non-null float64
1   pickup_datetime       187717 non-null object
2   pickup_longitude      187717 non-null float64
3   pickup_latitude       187717 non-null float64
4   dropoff_longitude     187716 non-null float64
5   dropoff_latitude      187715 non-null float64
6   passenger_count       187716 non-null float64
```

```
dtypes: float64(6), object(1)
```

```
df.describe() #To get statistics of each columns
```



	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
count	187717.000000	187717.000000	187717.000000	187716.000000	187715.000000	187716.000000
mean	11.357142	-72.527021	39.936938	-72.526449	39.923029	1.684209
std	9.879560	11.396782	7.823140	13.277703	6.850351	1.390584
min	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.985513	0.000000
25%	6.000000	-73.992070	40.734810	-73.991405	40.733835	1.000000
50%	8.500000	-73.981827	40.752611	-73.980087	40.753040	1.000000
75%	12.500000	-73.967137	40.767160	-73.963685	40.767993	2.000000
max	499.000000	57.418457	1644.421482	1153.572603	872.697628	208.000000

filling missing values

```
df.isnull().sum()
```

```
fare_amount      0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude 1
dropoff_latitude  2
passenger_count  1
dtype: int64
```

```
df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(),inplace = True)
```

```
df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(),inplace = True)
```

```
df.isnull().sum()
```

```
fare_amount      0
pickup_datetime  0
pickup_longitude  0
pickup_latitude  0
dropoff_longitude 0
dropoff_latitude 0
passenger_count  1
dtype: int64
```

```
df.dtypes
```

```
fare_amount      float64
pickup_datetime   object
pickup_longitude  float64
pickup_latitude   float64
dropoff_longitude float64
dropoff_latitude  float64
passenger_count   float64
dtype: object
```

Column pickup_datetime is in wrong format (Object). Convert it to DateTime Format

```
df.pickup_datetime = pd.to_datetime(df.pickup_datetime, errors='coerce')
```

```
df.dtypes
```

```
fare_amount      float64
pickup_datetime  datetime64[ns, UTC]
pickup_longitude  float64
pickup_latitude   float64
dropoff_longitude float64
dropoff_latitude  float64
```

passenger_count float64

To segregate each time of date and time

```
df= df.assign(hour = df.pickup_datetime.dt.hour,
              day= df.pickup_datetime.dt.day,
              month = df.pickup_datetime.dt.month,
              year = df.pickup_datetime.dt.year,
              dayofweek = df.pickup_datetime.dt.dayofweek)
```

```
df.head()
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.723217	1
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.750325	1
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.772647	1
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.803349	3
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.761247	5



```
# drop the column 'pickup_datetime' using drop()
# 'axis = 1' drops the specified column
```

```
df = df.drop('pickup_datetime',axis=1)
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month
0	7.5	-73.999817	40.738354	-73.999512	40.723217	1.0	19	7	
1	7.7	-73.994355	40.728225	-73.994710	40.750325	1.0	20	17	
2	12.9	-74.005043	40.740770	-73.962565	40.772647	1.0	21	24	
3	5.3	-73.976124	40.790844	-73.965316	40.803349	3.0	8	26	
4	16.0	-73.925023	40.744085	-73.973082	40.761247	5.0	17	28	



df.dtypes

```
fare_amount      float64
pickup_longitude  float64
pickup_latitude   float64
dropoff_longitude float64
dropoff_latitude  float64
passenger_count   float64
hour              int64
day              int64
month            int64
year             int64
dayofweek         int64
dtype: object
```

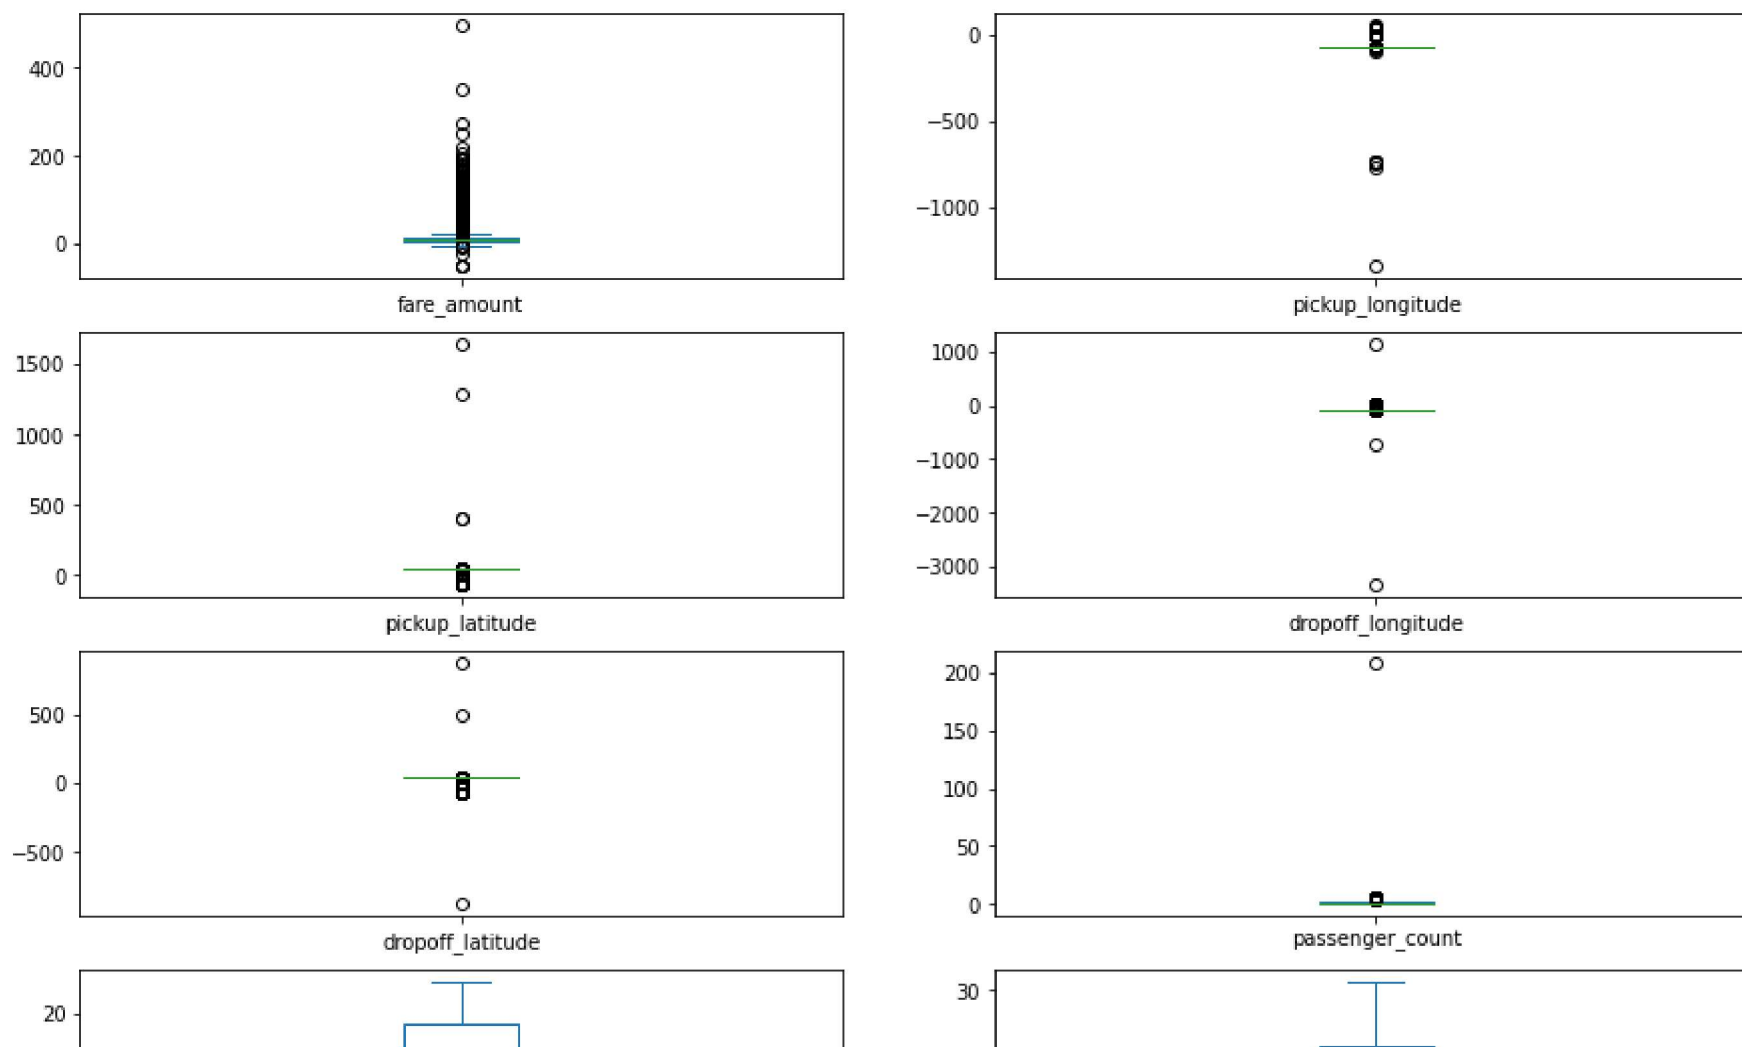
Checking outliers and filling them



```
df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20)) #Boxplot to check t
```

```

fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude  AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count  AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour             AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day              AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month            AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year             AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek        AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object

```





#Using the InterQuartile Range to fill the values

```
def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_whisker = Q1-1.5*IQR
    upper_whisker = Q3+1.5*IQR
    df[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
    return df1
```

```
def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df , c)
    return df1
```







```
df = treat_outliers_all(df , df.iloc[:, 0::])
```


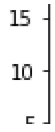
```
df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20)) #Boxplot shows that
```

```

fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude  AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count  AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour             AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day              AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month            AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year             AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek        AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object

```





```
!pip install haversine
import haversine as hs #Calculate the distance using Haversine to calculate the distance between to points. Can't use Euclai
travel_dist = []
for pos in range(len(df['pickup_longitude'])):
    long1,lati1,long2,lati2 = [df['pickup_longitude'][pos],df['pickup_latitude'][pos],df['dropoff_longitude'][pos],df['dr
    loc1=(lati1,long1)
    loc2=(lati2,long2)
    c = hs.haversine(loc1,loc2)
    travel_dist.append(c)

print(travel_dist)
df['dist_travel_km'] = travel_dist
df.head()
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting haversine

Downloading haversine-2.7.0-py2.py3-none-any.whl (6.9 kB)

Installing collected packages: haversine

Successfully installed haversine-2.7.0

IOPub data rate exceeded.

The notebook server will temporarily stop sending output

#Uber doesn't travel over 130 kms so minimize the distance

```
df= df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)]
```

```
print("Remaining observastions in the dataset:", df.shape)
```

Remaining observastions in the dataset: (187717, 12)

```
incorrect_coordinates = df.loc[(df.pickup_latitude > 90) | (df.pickup_latitude < -90) |
```

#Finding incorrect latitude (Less than or greater than 90) and longitude (greater than or less than 180)

```
incorrect_coordinates = df.loc[(df.pickup_latitude > 90) | (df.pickup_latitude < -90) |
                                (df.dropoff_latitude > 90) | (df.dropoff_latitude < -90) |
                                (df.pickup_longitude > 180) | (df.pickup_longitude < -180) |
                                (df.dropoff_longitude > 180) | (df.dropoff_longitude < -180)
                                ]
```

```
df.drop(incorrect_coordinates, inplace = True, errors = 'ignore')
```

```
df.head()
```

```

fare_amount pickup_longitude pickup_latitude dropoff_longitude dropoff_latitude passenger_count hour day month
df.isnull().sum()

```

```

fare_amount      0
pickup_longitude  0
pickup_latitude   0
dropoff_longitude 0
dropoff_latitude  0
passenger_count   1
hour              0
day              0
month            0
year             0
dayofweek         0
dist_travel_km    0
dtype: int64


```

```

sns.heatmap(df.isnull()) #Free for null values

```

<matplotlib.axes._subplots.AxesSubplot at 0x7febcb4f04d0>



```
corr = df.corr() #Function to find the correlation
```


```
corr
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
fare_amount	1.000000	0.153739	-0.110943	0.217658	-0.126442	0.015481
pickup_longitude	0.153739	1.000000	0.260639	0.425748	0.073404	-0.012587
pickup_latitude	-0.110943	0.260639	1.000000	0.050062	0.516141	-0.012590
dropoff_longitude	0.217658	0.425748	0.050062	1.000000	0.246565	-0.008615
dropoff_latitude	-0.126442	0.073404	0.516141	0.246565	1.000000	-0.006843
passenger_count	0.015481	-0.012587	-0.012590	-0.008615	-0.006843	1.000000
hour	-0.023112	0.010873	0.028834	-0.046539	0.019404	0.020272
day	0.005341	-0.003039	-0.001598	-0.004191	-0.003238	0.002537
month	0.030533	0.001178	0.001495	0.001733	-0.002644	0.010320
year	0.140274	0.009791	-0.012692	0.010938	-0.008649	-0.010034
dayofweek	0.012615	-0.024642	-0.042372	-0.003101	-0.031581	0.048305
dist_travel_km	0.785796	0.047197	-0.072991	0.154418	-0.053017	0.009810



```
fig,axis = plt.subplots(figsize = (10,6))
sns.heatmap(df.corr(),annot = True) #Correlation Heatmap (Light values means highly correlated)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7febc4d67450>



Dividing the dataset into feature and target values

```
x = df[['pickup_longitude','pickup_latitude','dropoff_longitude','dropoff_latitude','passenger_count','hour','day','month','year','dayofweek','dist_travel_km']]
```

```
y = df['fare_amount']
```

Dividing the dataset into training and testing dataset

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size = 0.33)
```

Linear Regression

```
from sklearn.linear_model import LinearRegression
regression = LinearRegression()
```

```
y_train = np.nan_to_num(y_train)
```

```
X_train= np.nan_to_num(X_train)
```

```
regression.fit(X_train,y_train)
```

```
LinearRegression()
```

```
regression.intercept_ #To find the linear intercept
```

```
3712.436030622846
```

```
regression.coef_ #To find the linear coefficient
```

```
array([ 2.59282080e+01, -7.29704623e+00,  2.02139949e+01, -1.80678549e+01,
        6.63557805e-02,  7.43757744e-03,  3.34792249e-03,  5.70327286e-02,
        3.67382492e-01, -3.65053829e-02,  1.85403830e+00])
```

```
prediction = regression.predict(X_test) #To predict the target values
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/base.py:444: UserWarning: X has feature names, but LinearRegression was
f"X has feature names, but {self.__class__.__name__} was fitted without"
```

```
print(prediction)
```

```
[10.24860732  6.74222458 10.52011102 ...  5.29117953 22.67254805
 8.45302783]
```



```
y_test
```

```
array([22.25,  5.7 ,  9.3 , ...,  4.5 , 22.25,  7.5 ])
```

Metrics Evaluation using R2, Mean Squared Error, Root Mean Squared Error

```
from sklearn.metrics import r2_score
r2_score(y_test, prediction)
```

```
0.6629487983734947
```

```
from sklearn.metrics import mean_squared_error
MSE = mean_squared_error(y_test, prediction)
MSE
```

```
9.935937305343638
```

```
RMSE = np.sqrt(MSE)
RMSE
```

```
3.152132183989694
```

Random Forest Regression

```
from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(n_estimators=100) #Here n_estimators means number of trees you want to build before making the prediction
rf.fit(X_train, y_train)
```

```
RandomForestRegressor()
```

```
y_pred = rf.predict(X_test)
y_pred
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/base.py:444: UserWarning: X has feature names, but RandomForestRegressor
```

```
f"X has feature names, but {self.__class__.__name__} was fitted without"
array([ 7.501 ,  6.275 , 10.443 , ...,  5.098 , 22.0455,  8.16  ])
```

Metrics evaluatin for Random Forest

```
R2_Random = r2_score(y_test,y_pred)
```

```
R2_Random
```

```
0.7937849498842736
```

```
MSE_Random = mean_squared_error(y_test,y_pred)
```

```
MSE_Random
```

```
6.079016480227936
```

```
RMSE_Random = np.sqrt(MSE_Random)
```

```
RMSE_Random
```

```
2.4655661581527144
```

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 3:44 PM

